


Introduction to High Performance Computing

Assignment 2

Wesley Liao
2019-09-27

1. Because the `dot_prod` is assigned to its current value plus the calculated vector product, this creates an opportunity for threads to use an old value of `dot_prod`. Code on page 2.
2. The order of the chunk execution and the printout of “Thread id: x working on index i” may change but the end result of the program will be the same.
3.
 - (a) For 3 OpenMP threads, the initialization for loop will run in all 3 threads, and each omp section will execute in parallel.
 - (b) For 5 OpenMP threads, the initialization for loop will run in all 5 threads, but after that there are only 3 sections, so there will be only work for 3 of the threads.
 - (c) With just 1 OpenMP thread, each of the three sections will essentially be executed sequentially.
 - (d) In general, there will be very little to no performance improvement when there are more threads than parallel sections.
4. Code on page 3. When compiled, `p4` accepts 1 argument, an integer number of threads to run. Performance was measured using the “time” utility, running 1 through 8 threads 4 times each and plotting the resulting time reported below:



Code for part 1:

```
#include <stdio.h>
#include <omp.h>

// compute the dot product of two vectors

int main() {
    int const N=100;
    int i;
    double a[N], b[N];
    double dot_prod = 0.0;

    int thread_id;
    //Arbitrarily initialize vectors a and b
    for (i = 0; i < N; i++) {
        a[i] = 3.14;
        b[i] = 6.67;
    }

    #pragma omp parallel private(thread_id)
    {
        thread_id = omp_get_thread_num();
        printf("This thread is: %d\n", thread_id);
        #pragma omp for
        for(i=0; i < N; i++) {
            // sum up the element-wise product of the two arrays
            #pragma omp atomic
            {
                dot_prod = dot_prod + (a[i] * b[i]);
            }
        }
    }

    printf("Dot product of the two vectors is %g\n", dot_prod);

    return 0;
}
```

Code for part 4:

```
#include <stdio.h>
#include <stdlib.h>
// link with -lm at compile time
#include <math.h>
#include <omp.h>

int main(int argc, char* argv[]) {

    int threads = atoi(argv[1]);
    int const N=1000;
    int i, j, ij;

    double A[N*N];
    double x[N], b[N];

    // initialize the matrix and the vector
    #pragma omp parallel for num_threads(threads)
    for (ij = 0; ij < N*N; ij++) {
        A[ij] = sin(0.01*(ij));
    }

    #pragma omp parallel for num_threads(threads)
    for (i = 0; i < N; i++) {
        b[i] = cos(0.01*i);
        x[i] = 0.0;
    }

    // matrix vector multiplication
    #pragma omp parallel for num_threads(threads)
    for (ij = 0; ij < N*N; ij++) {
        #pragma omp atomic
        x[(ij - (ij%N))/N] = x[(ij - (ij%N))/N] + (A[ij]*b[ij%N]);
    }

    printf("x[%d] = %g\n", 505, x[505]);

    return 0;
}
```